

**Exercise 1:**

**1) AI vs ML vs DL**

- *Main idea:* Artificial Intelligence is the broad field of building systems that perform tasks associated with intelligence; Machine Learning is a subfield that learns from data to improve task performance; Deep Learning is a subfield of ML based on multi-layer neural networks.
- ML and statistics overlap strongly in methods and goals; differences are often historical and terminological.
- *Lecturer comment:* Keep the taxonomy clear to avoid using AI as a catch-all. Position ML as data-driven prediction and DL as a flexible function class inside ML. Translate terms across ML and statistics to reduce confusion.

**2) Learning paradigms: supervised, unsupervised, reinforcement learning**

- *Main idea:* Supervised learning uses labeled pairs  $(x, y)$  to learn  $x \rightarrow y$ . Unsupervised learning uses only  $x$  to discover structure (clustering, dimensionality reduction, anomaly detection). Reinforcement learning learns via interaction and rewards, optimizing long-run return.
- *Lecturer comment:* Specify underlying supervised task (regression, binary or multiclass classification) before picking a learner. Note that unsupervised learning is not a single objective (can be clustering, dimensionality reduction, anomaly detection optimize different criteria, ...); RL rewards can be sparse, delayed, and noisy.

**3) Data in supervised learning: features vs target; labeled vs unlabeled; Iris example**

- *Main idea:* The target  $y$  is the prediction goal; features  $x$  describe objects. We assume a predictive relation between  $x$  and  $y$ .
- Distinguish labeled from unlabeled data. The Iris dataset (150 samples, 4 measurements) is a clean toy example for classification.
- *Lecturer comment:* Iris is pedagogically useful but unusually clean; do not generalize its simplicity to real data. When first seeing any dataset, label the data type: supervised with  $(x, y)$  or unsupervised with only  $x$ .

**4) Data types and encodings: numeric vs categorical; one-hot vs dummy; ordinal encodings**

- *Main idea:* Many learners expect numeric inputs, so categorical variables need encoding.
- One-hot uses  $k$  indicator columns; dummy uses  $k - 1$  to avoid a redundant column. Ordinal encodings are appropriate only when a natural order exists.
- *Lecturer comment:* Decision trees can handle categoricals more directly; linear models typically require encodings. Use dummy coding to avoid singular design matrices; avoid imposing false order with naive integer encoding (where categories are assigned numbers).

**5) Data-generating process and the i.i.d. assumption**

- *Main idea:* We assume data arise from an unknown distribution  $P_{X,Y}$ . Standard analyses treat samples as independent and identically distributed.
- *Lecturer comment:* I.i.d. is a modeling convenience. Time series, networks, and nonstationary streams often violate it. Keep the assumption now, but flag where it can fail.

**6) Tasks: regression vs classification**

- *Main idea:* If the target is numeric, the task is regression; if the target is categorical with  $g$  classes, the task is classification.

- *Lecturer comment:* Tie the task directly to target type to avoid evaluation mismatches later. The same domain can be framed for prediction or explanation; be explicit which you pursue.

## 7) Predict vs explain

- *Main idea:* Predict focuses on out-of-sample accuracy; explain focuses on understanding patterns and relations. Both require adequate fit.
- *Lecturer comment:* Explanation here is descriptive, not necessarily causal. Stakeholders typically want both a strong predictor and an understandable story; decide which objective dominates before modeling.

## 8) Models and hypothesis spaces

- *Main idea:* A model  $f$  maps features to outputs; in classification  $f$  often produces scores or probabilities. We restrict to a hypothesis space  $H$  to encode structure and make learning feasible.
- *Lecturer comment:* Demystify models as functions  $f : \mathcal{X} \rightarrow \mathcal{Y}$ . The choice of  $H$  (linear functions, trees, neural networks, etc.) constrains what can be learned.

## 9) Parameters, identifiability, and illustrative examples

- *Main idea:* Models in  $H$  share a parametric form with parameter vector  $\theta$ ; choosing  $\theta$  fixes  $f$ . Some classes are non-identifiable (distinct  $\theta$  yield the same  $f$ ).
- Examples: univariate linear  $f(x) = \theta_0 + \theta_1 x$ , bivariate quadratics, and RBF networks with centers, widths, and weights. Hyperparameters such as number of centers  $k$  or bandwidth define families prior to training.
- *Lecturer comment:* Parameterization operationalizes  $H$ . Non-identifiability matters for optimization and interpretation. Hyperparameters set the scope of the family before fitting.

## 10) Learners (inducers)

- *Main idea:* A learner receives training data  $D$  and controls  $\Lambda$ , and selects  $f \in H$  (or  $\theta$ ) that minimizes empirical risk on  $D$ .
- Formal mapping:  $\mathcal{I} : \mathbb{D} \times \Lambda \rightarrow H$ .
- *Lecturer comment:* Separate the algorithm from the model. Any training code instantiates the mapping from data and controls to a fitted function  $f$ .

## 11) Loss functions

- *Main idea:* The loss  $L(y, f(x))$  measures pointwise error. For regression,  $L1$  and  $L2$  are common; classification uses appropriate losses or surrogates. Aggregating losses gives empirical risk.
- *Lecturer comment:*  $L2$  magnifies large residuals;  $L1$  is more robust. Keep a tiny numeric example handy to make robustness concrete.

## 12) Risk minimization: theoretical vs empirical risk

- *Main idea:* Theoretical risk  $R(f) = E[L(Y, f(X))]$  is defined under the unknown data distribution. We approximate it by empirical risk  $R_{\text{emp}}(f) = \frac{1}{n} \sum_{i=1}^n L(y_i, f(x_i))$  on i.i.d. samples.
- *Lecturer comment:* Alternatives exist (estimate  $P_{X,Y}$ , or assume a parametric form), but ERM is the default in this chapter because  $P_{X,Y}$  is unknown.

## 13) Empirical Risk Minimization (ERM)

- *Main idea:* ERM chooses  $f = \arg \min_{f \in H} R_{\text{emp}}(f)$ , equivalently  $\theta = \arg \min_{\theta} R_{\text{emp}}(\theta)$ .
- For finite  $H$  one could tabulate risks; for infinite  $H$  we optimize a surface over parameters.
- *Lecturer comment:* Reducing learning to numerical optimization is powerful but incomplete; generalization and model selection enter next.

## 14) Optimization for learning

- *Main idea:* Many learners solve  $\min_{\theta} R_{\text{emp}}(\theta)$  via numerical optimization. Gradient descent updates parameters along negative gradients. The learning rate controls step size.
- Local minima vs global minima: in practice, good local optima often suffice. Some models admit analytic solutions (for example, OLS in linear regression under standard conditions).

- *Lecturer comment:* Trade off speed and stability with the step size. Name families briefly: first order (GD, SGD), second order (Newton type), and practical variants (momentum, Adam).

## 15) Components of supervised learning

- *Main idea:* Many supervised learners decompose into hypothesis space  $H$  plus loss/risk plus optimization. Regularization can be folded into the risk.
- Examples of choices:  $H$  (linear, trees, nets), losses (MSE, NLL, misclassification surrogates), optimizers (analytical, gradient based, combinatorial).
- *Lecturer comment:* Keep returning to the triad  $H + \text{risk} + \text{optimization}$  as an organizing lens to compare algorithms. Regularization integrates naturally into the objective.