

Exercise 3 – Classification I

Introduction to Machine Learning

Exercise 1: Logistic regression – softmax link and decision boundary

The whole exercise is only for lecture group A!

Learning goals

1. Solve “show equivalence”-type questions
2. Derive the decision boundary of the logistic classifier

Binary logistic regression is a special case of multiclass logistic, or softmax, regression. The softmax function is the multiclass analogue to the logistic function, transforming scores $\theta^\top \mathbf{x}$ to values in the range $[0, 1]$ that sum to one. The softmax function is defined as:

$$\pi_k(\mathbf{x}|\theta) = \frac{\exp(\theta_k^\top \mathbf{x})}{\sum_{j=1}^g \exp(\theta_j^\top \mathbf{x})}, k \in \{1, \dots, g\}$$

- a) Show that logistic and softmax regression are equivalent for $g = 2$.
- b) Recall that in the binary case we predict $\hat{y} = 1$ iff $\pi(\mathbf{x}) \geq \alpha$ for a threshold $\alpha \in (0, 1)$. Show that the resulting decision boundary is a (linear) hyperplane.

Hint

Derive the value of $\theta^\top \mathbf{x}$ (depending on α) starting from which you predict $\hat{y} = 1$ rather than $\hat{y} = 0$.

Exercise 2: Hyperplanes

Learning goals

1. Understand that hyperplanes bisect the space with a linear boundary
2. Get a feeling for coefficients in hyperplane equations
3. See how the classification threshold α shifts the logistic decision boundary

Linear classifiers like logistic regression learn a decision boundary that takes the form of a (linear) hyperplane. Hyperplanes are defined by equations $\theta^\top \mathbf{x} = b$ with coefficients θ and a scalar $b \in \mathbb{R}$.

- a) In order to see that such expressions actually describe hyperplanes, consider $\theta^\top \mathbf{x} = \theta_0 + \theta_1 x_1 + \theta_2 x_2 = 0$. Sketch the hyperplanes given by the following coefficients and explain the difference between the parameterizations:

- $\theta_0 = 0, \theta_1 = \theta_2 = 1$
- $\theta_0 = 1, \theta_1 = \theta_2 = 1$
- $\theta_0 = 0, \theta_1 = 1, \theta_2 = 2$

- b) Beyond its orientation, a linear classifier also has an *offset* that determines where exactly the boundary sits. Recall that a logistic regression model estimates $\pi(\mathbf{x}) = \frac{1}{1 + \exp(-\theta^\top \mathbf{x})}$ and predicts $\hat{y} = 1$ iff $\pi(\mathbf{x}) \geq \alpha$ for a threshold $\alpha \in (0, 1)$. Its decision boundary is therefore the hyperplane

$$\theta^\top \mathbf{x} = -\log\left(\frac{1}{\alpha} - 1\right) =: b(\alpha).$$

Consider the fixed model $\theta = (\theta_0, \theta_1, \theta_2)^\top = (0, 1, 1)^\top$. Compute the offset $b(\alpha)$ for $\alpha \in \{0.25, 0.5, 0.75\}$, sketch the three resulting decision boundaries in a single (x_1, x_2) plot, and describe how the boundary – and the region predicted as class 1 – changes as α grows. Does the orientation of the boundary change?

Exercise 3: Decision Boundaries & Thresholds in Logistic Regression

Learning goals

1. Understand that logistic regression finds a linear decision boundary
2. Get a feeling for how the parameters θ shape predicted probabilities and model confidence
3. Understand the role of the decision threshold α

Recall from Exercise 2 that a logistic regression model estimates the probability $p(y = 1|\mathbf{x}, \theta) = \pi(\mathbf{x}|\theta)$ and assigns $\hat{y} = 1$ iff $\pi(\mathbf{x}|\theta) \geq \alpha$ for a threshold $\alpha \in (0, 1)$. Whereas Exercise 2 b) varied the threshold for a fixed model, we now study how the parameters θ themselves shape the predicted probabilities.

- a) Below you see the logistic function for a binary classification problem with two input features for different values $\theta^\top = (\theta_1, \theta_2)^\top$ (plots 1-3) as well as α (plot 4). What can you deduce for the values of θ_1 , θ_2 , and α ? What are the implications for classification in the different scenarios? In particular, which of the plots (1)–(3) shows the most confident model?
- b) In Exercise 2 b) we simply read off the offset $b(0.5) = 0$ for the common threshold $\alpha = 0.5$. Derive this special case directly, starting from $\pi(\mathbf{x}) = 0.5$.
- c) Explain when it might be sensible to set α to 0.5.

Exercise 4: Thresholds and predictions in practice

The whole exercise is only for lecture group B!

Learning goals

1. Compute predictions of a logistic model by hand
2. Understand how the threshold α trades off the two error types


Consider a fitted binary logistic regression model with parameter vector $\theta = (\theta_0, \theta_1, \theta_2)^\top = (-1, 2, 1)^\top$, so that

$$\pi(\mathbf{x}) = \frac{1}{1 + \exp(-(\theta_0 + \theta_1 x_1 + \theta_2 x_2))}.$$


Recall that we predict $\hat{y} = 1$ iff $\pi(\mathbf{x}) \geq \alpha$.

- a) Compute the predicted probability $\pi(\mathbf{x})$ for the following observations and classify them using the default threshold $\alpha = 0.5$:
 - $\mathbf{x}^{(1)} = (x_1, x_2)^\top = (1, 1)^\top$
 - $\mathbf{x}^{(2)} = (1, -1)^\top$
 - $\mathbf{x}^{(3)} = (0, 0)^\top$
- b) For observation $\mathbf{x}^{(1)}$, from which threshold α onwards would the model predict $\hat{y} = 0$ instead?


Plot (1)




Plot (2)



Plot (3)



Plot (4)



- c) The model is used in a hospital to decide whether to start an expensive and invasive therapy ($\hat{y} = 1$). Would you rather choose a high or a low threshold α ? Which type of error (false positive / false negative) does your choice avoid, and at what cost?