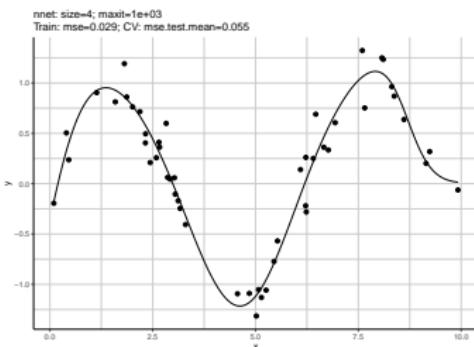


Introduction to Machine Learning

Neural Networks

Universal Approximation



Learning goals

- Universal approximation theorem for one-hidden-layer neural networks
- The pros and cons of a low approximation error

UNIVERSAL APPROXIMATION PROPERTY

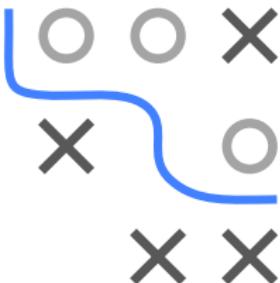
Theorem. Let $\sigma : \mathbb{R} \rightarrow \mathbb{R}$ be a continuous, non-constant, bounded, and monotonically increasing function. Let $C \subset \mathbb{R}^p$ be compact, and let $\mathcal{C}(C)$ denote the space of continuous functions $C \rightarrow \mathbb{R}$. Then, given a function $g \in \mathcal{C}(C)$ and an accuracy $\varepsilon > 0$, there exists a hidden layer size $m \in \mathbb{N}$ and a set of coefficients $W_j \in \mathbb{R}^p$, $u_j, b_j \in \mathbb{R}$ (for $j \in \{1, \dots, m\}$), such that

$$f : C \rightarrow \mathbb{R}; \quad f(\mathbf{x}) = \sum_{j=1}^m u_j \cdot \sigma(W_j^T \mathbf{x} + b_j)$$

is an ε -approximation of g , that is,

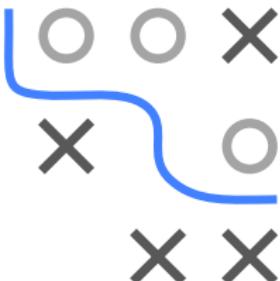
$$\|f - g\|_\infty := \max_{x \in C} |f(x) - g(x)| < \varepsilon .$$

The theorem extends trivially to multiple outputs.



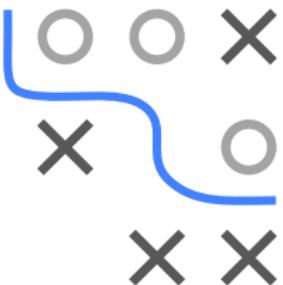
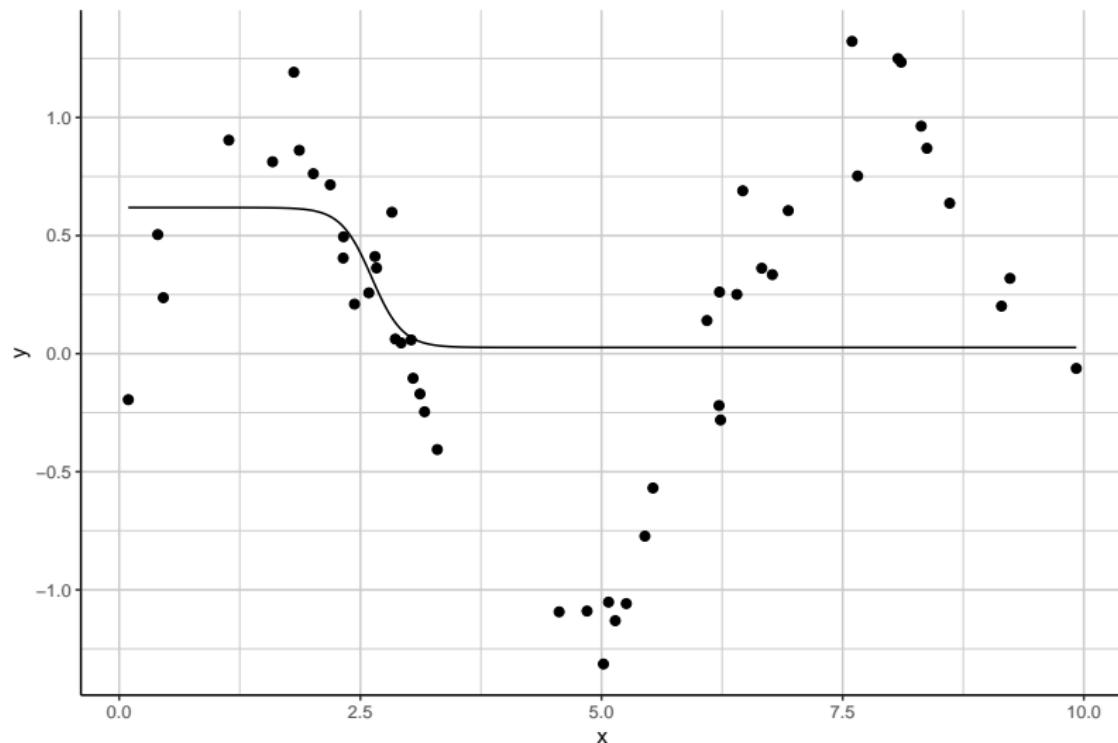
EXAMPLE : REGRESSION/CLASSIFICATION

- Let's look at a few examples of the types of functions and decisions boundaries learnt by neural networks (with a **single** hidden layer) of various sizes.
- "size" here refers to the number of neurons in the hidden layer.
- The number of "iterations" in the following slides corresponds to the number of steps of the applied iterative optimization algorithm (stochastic gradient descent).



REGRESSION EX.: 1000 TRAINING ITERATIONS

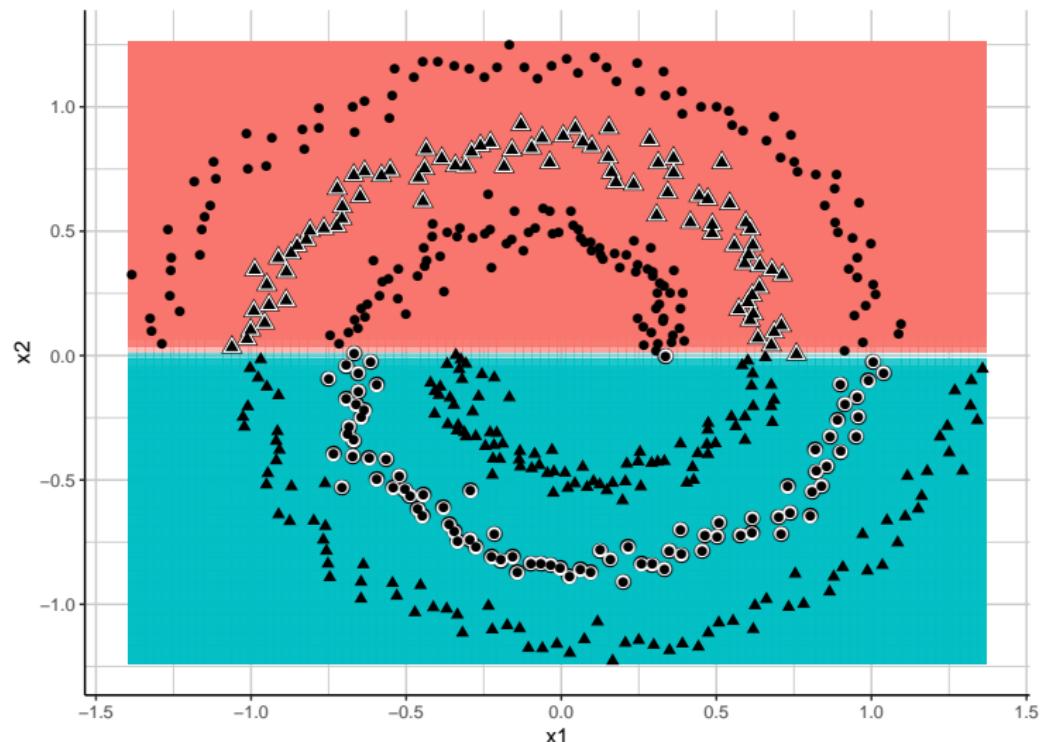
nnet: size=1; maxit=1e+03
Train: mse=0.391; CV: mse.test.mean=0.419



CLASSIFICATION: 500 TRAINING ITERATIONS

nnet: size=1; maxit=500

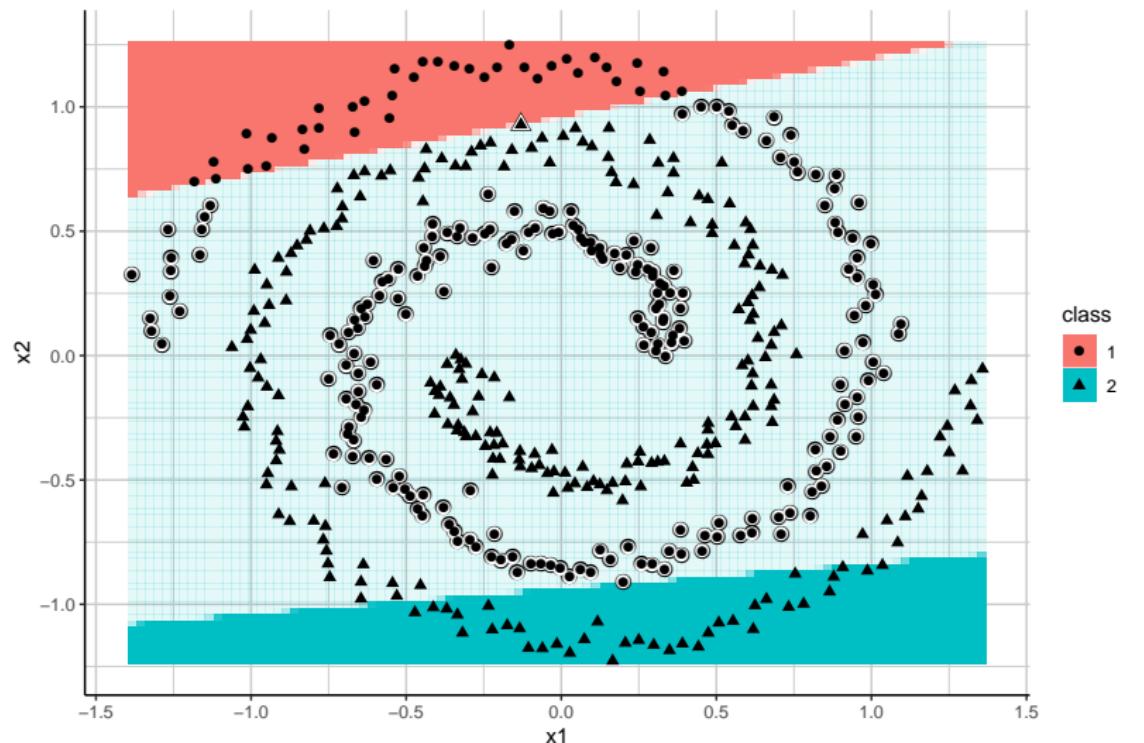
Train: mmce=0.336; CV: mmce.test.mean=0.346



CLASSIFICATION: 500 TRAINING ITERATIONS

nnet: size=2; maxit=500

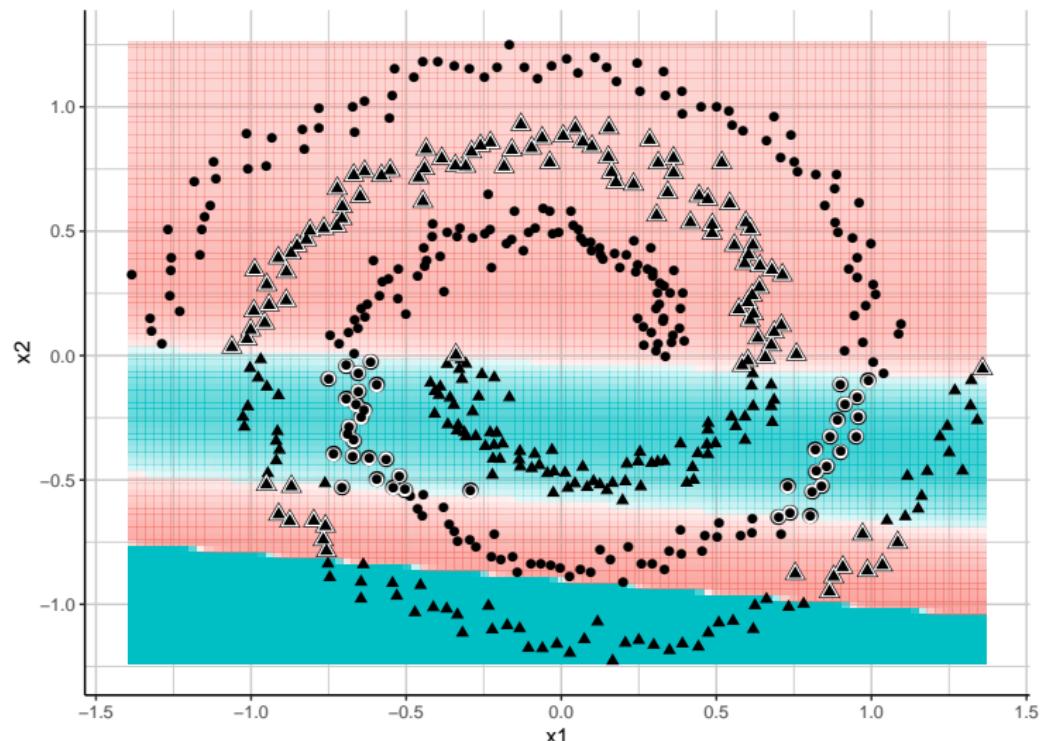
Train: mmce=0.426; CV: mmce.test.mean=0.412



CLASSIFICATION: 500 TRAINING ITERATIONS

nnet: size=3; maxit=500

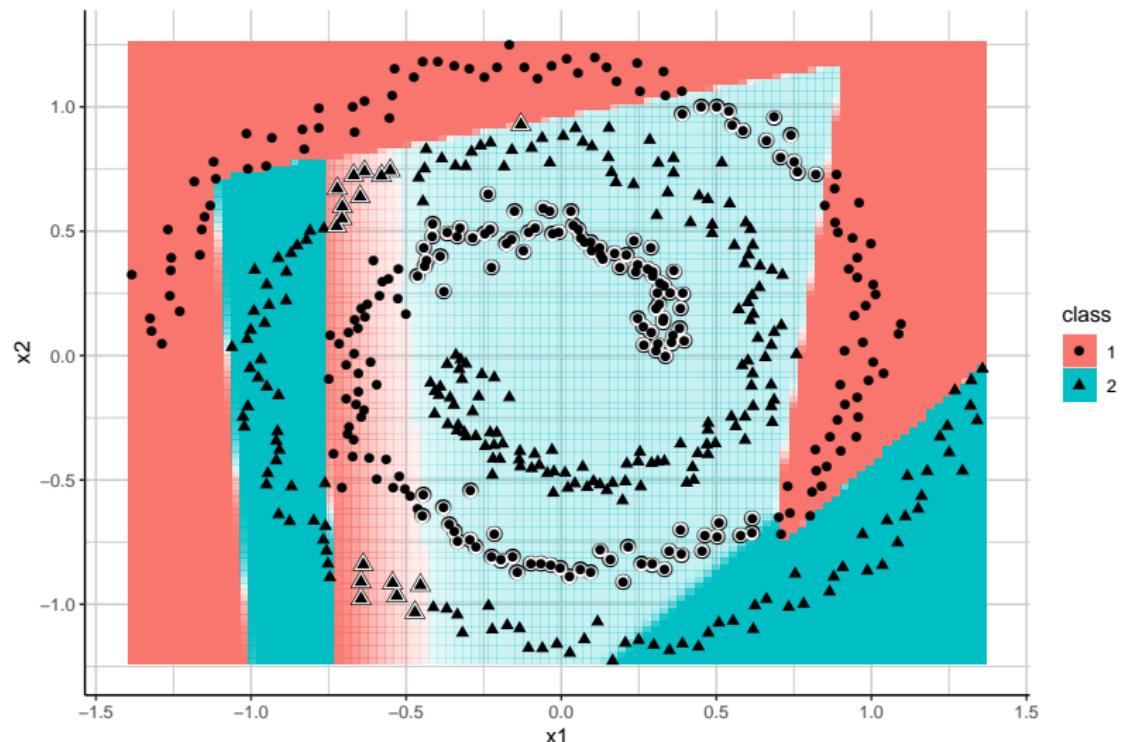
Train: mmce=0.290; CV: mmce.test.mean=0.374



CLASSIFICATION: 500 TRAINING ITERATIONS

nnet: size=5; maxit=500

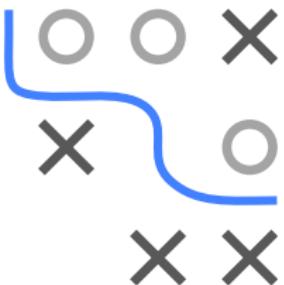
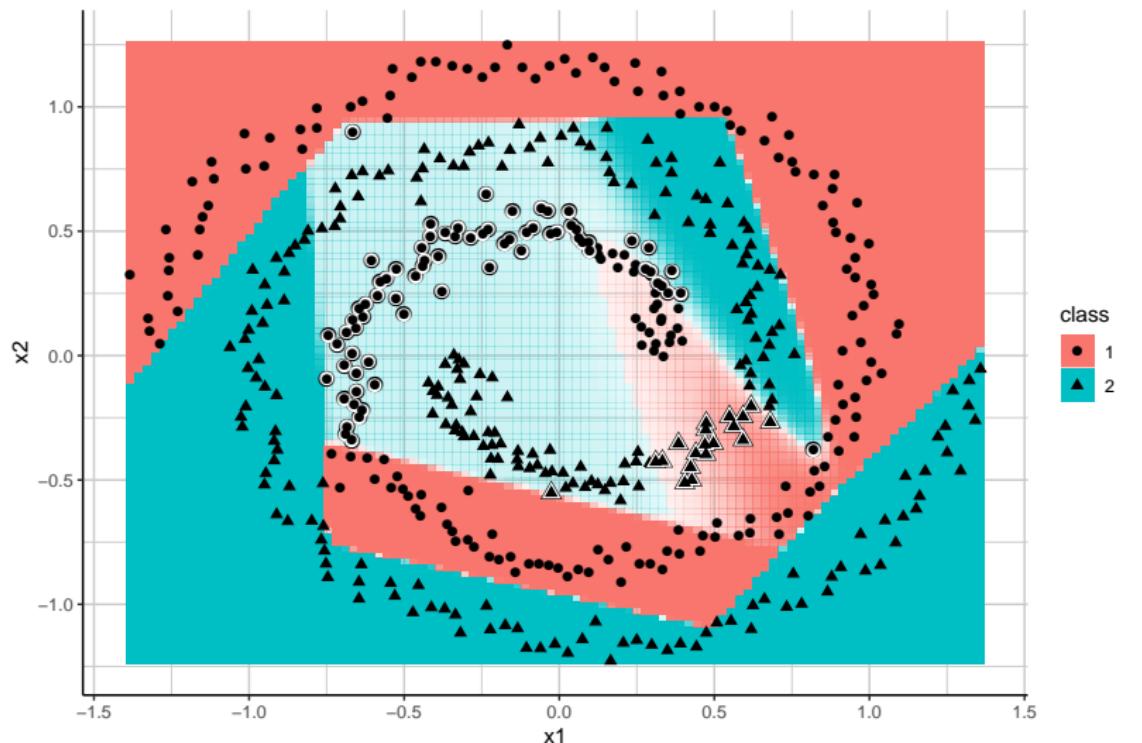
Train: mmce=0.272; CV: mmce.test.mean=0.322



CLASSIFICATION: 500 TRAINING ITERATIONS

nnet: size=10; maxit=500

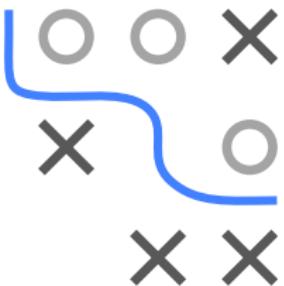
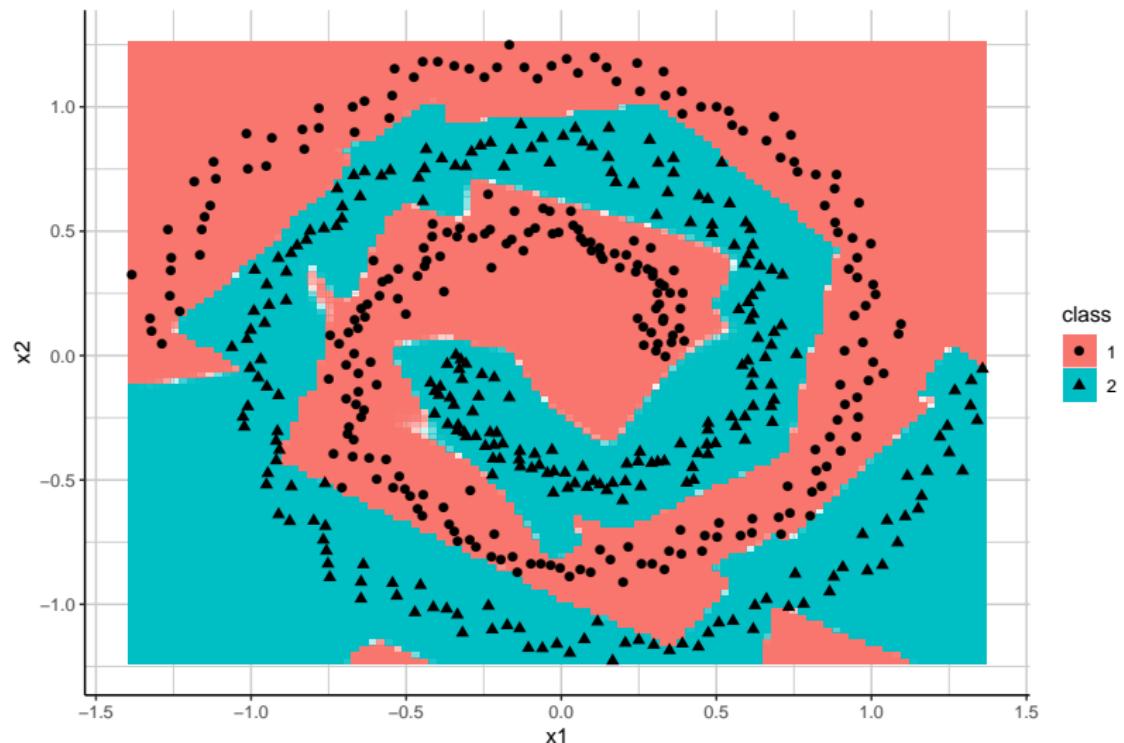
Train: mmce=0.184; CV: mmce.test.mean=0.106



CLASSIFICATION: 500 TRAINING ITERATIONS

nnet: size=30; maxit=500

Train: mmce=0.000; CV: mmce.test.mean=0.034



CLASSIFICATION: 500 TRAINING ITERATIONS

nnet: size=50; maxit=500

Train: mmce=0.000; CV: mmce.test.mean=0.026

